

Lab 6 Report – Abdul Hannan Khan

Task 1: Wine Quality Classification

Explanation:

This program classifies three types of wine chemical features. It trains a K-Nearest Neighbors (KNN) machine learning model with k=5 to predict wine categories. The model achieves about 74% accuracy on test data and creates scatter plots showing how different chemical features separate the wine classes. Finally, it displays the first 10 predictions versus actual values to show how well the model performs.


Code:

```
👤 AbdulHannanKhan_Task1_KNNClassification_WineQuality.py > ...
1  import numpy as np
2  import matplotlib.pyplot as plt
3  from sklearn.datasets import load_wine
4  from sklearn.model_selection import train_test_split
5  from sklearn.neighbors import KNeighborsClassifier
6  from sklearn.metrics import accuracy_score
7
8  wine = load_wine()
9  X = wine.data
10 y = wine.target
11
12 print("Dataset loaded successfully!")
13 print(f"Number of wine samples: {X.shape[0]}")
14 print(f"Number of features: {X.shape[1]}")
15 print(f"Wine types: {np.unique(y)}")
16 print(f"Feature names: {wine.feature_names}")
17
18 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
19
20 print(f"\nTraining set size: {X_train.shape[0]} wines")
21 print(f"Testing set size: {X_test.shape[0]} wines")
22
23 knn = KNeighborsClassifier(n_neighbors=5)
24 knn.fit(X_train, y_train)
25
26 print("\nKNN model trained with k=5")
27
28 y_pred = knn.predict(X_test)
29 accuracy = accuracy_score(y_test, y_pred)
```

```
50
31 print(f"Model accuracy on test set: {accuracy:.2f} ({accuracy*100:.1f}%)")
32
33 plt.figure(figsize=(12, 5))
34
35 plt.subplot(1, 2, 1)
36 colors = ['red', 'green', 'blue']
37 for i, color in enumerate(colors):
38     plt.scatter(X[y == i, 0], X[y == i, 1], c=color, label=f'Wine Class {i}', alpha=0.7)
39
40 plt.xlabel('Alcohol Content')
41 plt.ylabel('Malic Acid')
42 plt.title('Wine Classes: Alcohol vs Malic Acid')
43 plt.legend()
44 plt.grid(True, alpha=0.3)
45
46 plt.subplot(1, 2, 2)
47 for i, color in enumerate(colors):
48     plt.scatter(X[y == i, 0], X[y == i, 9], c=color, label=f'Wine Class {i}', alpha=0.7)
49
50 plt.xlabel('Alcohol Content')
51 plt.ylabel('Color Intensity')
52 plt.title('Wine Classes: Alcohol vs Color Intensity')
53 plt.legend()
54 plt.grid(True, alpha=0.3)
55
56 plt.tight_layout()
57 plt.show()
58
59 print("\nFirst 10 predictions vs actual:")
60 print("Predicted | Actual")
61 print("-" * 18)
62 for i in range(10):
63     print(f"    {y_pred[i]} | {y_test[i]}")
```

Output:

Figure 1



```
Dataset loaded successfully!
Number of wine samples: 178
Number of features: 13
Wine types: [0 1 2]
Feature names: ['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash', 'magnesium', 'total_phenols', 'flavanoids', 'nonflavanoid_phenols', 'proanthocyanins', 'color_intensity', 'hue', 'od280/od315_of_diluted_wines', 'proline']

Training set size: 124 wines
Testing set size: 54 wines

KNN model trained with k=5
Model accuracy on test set: 0.74 (74.1%)

First 10 predictions vs actual:
Predicted | Actual
-----|-----
2 | 0
0 | 0
2 | 2
0 | 0
1 | 1
0 | 0
2 | 1
2 | 2
1 | 1
0 | 2
```

Task 2: Breast Cancer Detection

Explanation:

This program demonstrates how features scaling improves KNN performance for breast cancer detection. It compares two models – one using original unscaled features and another using standardized features. The results show that scaling significantly boosts accuracy because KNN relies on distance calculations, and scaling ensures all features contribute equally. The visualization clearly demonstrates the performance improvement achieved through proper data preprocessing.

Code:

```
AbdulHannanKhan_Task2_KNNClassification_BreastCancerDetection.py > ...
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.datasets import load_breast_cancer
4 from sklearn.model_selection import train_test_split
5 from sklearn.neighbors import KNeighborsClassifier
6 from sklearn.metrics import accuracy_score, classification_report
7 from sklearn.preprocessing import StandardScaler
8
9 print("==> Breast Cancer Detection with KNN ==\n")
10
11 cancer = load_breast_cancer()
12 X = cancer.data
13 y = cancer.target
14
15 print("Dataset loaded successfully!")
16 print(f"Number of samples: {X.shape[0]}")
17 print(f"Number of features: {X.shape[1]}")
18 print(f"Target distribution:")
19 print(f"- Malignant (0): {np.sum(y == 0)} samples")
20 print(f"- Benign (1): {np.sum(y == 1)} samples")
21
22 print("\nFirst 5 features and their ranges (min to max):")
23 for i in range(5):
24     feature_name = cancer.feature_names[i]
25     min_val = np.min(X[:, i])
26     max_val = np.max(X[:, i])
27     print(f"{feature_name:25}: {min_val:8.2f} to {max_val:8.2f}")
28
29 print("\n" + "="*50)
30 print("BASELINE MODEL (Without Feature Scaling)")
31 print("-"*50)
32
33 X_train_unscaled, X_test_unscaled, y_train, y_test = train_test_split(
34     X, y, test_size=0.2, random_state=42, stratify=y
35 )
36
```

```

37 print(f"Training set: {X_train_unscaled.shape[0]} samples")
38 print(f"Testing set: {X_test_unscaled.shape[0]} samples")
39
40 knn_unscaled = KNeighborsClassifier(n_neighbors=7)
41 knn_unscaled.fit(X_train_unscaled, y_train)
42
43 y_pred_unscaled = knn_unscaled.predict(X_test_unscaled)
44 accuracy_unscaled = accuracy_score(y_test, y_pred_unscaled)
45
46 print(f"\nAccuracy without scaling: {accuracy_unscaled:.4f} ({accuracy_unscaled*100:.2f}%)")
47
48 print("\n" + "="*50)
49 print("IMPROVED MODEL (With Feature Scaling)")
50 print("="*50)
51
52 scaler = StandardScaler()
53 X_scaled = scaler.fit_transform(X)
54
55 print("Features scaled using StandardScaler")
56
57 print("\nExample of scaling effect on first feature:")
58 print(f"Before scaling - Mean: {np.mean(X[:, 0]):.2f}, Std: {np.std(X[:, 0]):.2f}")
59 print(f"After scaling - Mean: {np.mean(X_scaled[:, 0]):.2f}, Std: {np.std(X_scaled[:, 0]):.2f}")
60
61 X_train_scaled, X_test_scaled, y_train, y_test = train_test_split(
62     X_scaled, y, test_size=0.2, random_state=42, stratify=y
63 )
64
65 knn_scaled = KNeighborsClassifier(n_neighbors=7)
66 knn_scaled.fit(X_train_scaled, y_train)
67
68 y_pred_scaled = knn_scaled.predict(X_test_scaled)
69 accuracy_scaled = accuracy_score(y_test, y_pred_scaled)

```

```

70     print(f"\nAccuracy with scaling: {accuracy_scaled:.4f} ({accuracy_scaled*100:.2f}%}")
71
72     print("\n" + "="*50)
73     print("COMPARISON RESULTS")
74     print("="*50)
75
76     improvement = accuracy_scaled - accuracy_unscaled
77     improvement_percent = (improvement / accuracy_unscaled) * 100
78
79     print(f"Baseline accuracy (unscaled): {accuracy_unscaled*100:.2f}%")
80     print(f"Improved accuracy (scaled): {accuracy_scaled*100:.2f}%")
81     print(f"Improvement: +{improvement*100:.2f}% points")
82     print(f"Relative improvement: +{improvement_percent:.1f}%")
83
84     plt.figure(figsize=(12, 5))
85
86     plt.subplot(1, 2, 1)
87     models = ['Without Scaling', 'With Scaling']
88     accuracies = [accuracy_unscaled, accuracy_scaled]
89     colors = ['red', 'green']
90
91     bars = plt.bar(models, accuracies, color=colors, alpha=0.7)
92     plt.ylabel('Accuracy Score')
93     plt.title('KNN Performance: Scaled vs Unscaled Features')
94     plt.ylim(0, 1.0)
95
96     for bar, accuracy in zip(bars, accuracies):
97         plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.01,
98                  f'{accuracy:.3f}', ha='center', va='bottom')
99
100    plt.subplot(1, 2, 2)
101    features_to_show = 3
102    x_pos = np.arange(features_to_show)
103
104

```

```

105 std_before = [np.std(X[:, i]) for i in range(features_to_show)]
106 std_after = [np.std(X_scaled[:, i]) for i in range(features_to_show)]
107
108 plt.bar(x_pos - 0.2, std_before, 0.4, label='Before Scaling', alpha=0.7, color='red')
109 plt.bar(x_pos + 0.2, std_after, 0.4, label='After Scaling', alpha=0.7, color='green')
110
111 plt.xlabel('Feature Index')
112 plt.ylabel('Standard Deviation')
113 plt.title('Feature Scale Comparison')
114 plt.xticks(x_pos, [f'Feature {i+1}' for i in range(features_to_show)])
115 plt.legend()
116 plt.grid(True, alpha=0.3)
117
118 plt.tight_layout()
119 plt.show()
120
121 print("\n" + "="*50)
122 print("DETAILED CLASSIFICATION REPORT (With Scaling)")
123 print("="*50)
124 print(classification_report(y_test, y_pred_scaled,
125 | | | | | target_names=['Malignant', 'Benign']))
126
127 print("\n" + "="*50)
128 print("WHY FEATURE SCALING IS CRITICAL FOR KNN")
129 print("="*50)
130
131 print(f"""
132 K-Nearest Neighbors (KNN) is a distance-based algorithm that calculates
133 the distance between data points to make predictions.
134

```


```

135 WHY SCALING MATTERS:
136 1. Features with larger numerical ranges dominate distance calculations
137     Example:
138     - Feature A range: 0-1 (small influence)
139     - Feature B range: 0-1000 (huge influence)
140     Without scaling, Feature B would completely dominate the distance!
141
142 2. Equal Contribution:
143     - Scaling ensures all features contribute equally to the distance calculation
144     - No single feature unfairly influences the results
145
146 3. Better Performance:
147     - As shown above, scaling improved accuracy from {accuracy_unscaled*100:.1f}% to {accuracy_scaled*100:.1f}%
148     - This is a {improvement_percent:.1f}% relative improvement!
149
150 StandardScaler transforms features to have:
151 - Mean = 0
152 - Standard Deviation = 1
153 This puts all features on the same scale without changing their distribution shape.
154 """

```

Output:

Figure 1



Home | Back | Forward | Refresh | Search | Help

```
Dataset loaded successfully!
Number of samples: 569
Number of features: 30
Target distribution:
- Malignant (0): 212 samples
● - Benign (1): 357 samples

First 5 features and their ranges (min to max):
mean radius      :    6.98 to    28.11
mean texture     :    9.71 to    39.28
mean perimeter   :   43.79 to   188.50
mean area        :  143.50 to 2501.00
mean smoothness  :    0.05 to    0.16

=====
===== BASELINE MODEL (Without Feature Scaling) =====
=====
Training set: 455 samples
Testing set: 114 samples

Accuracy without scaling: 0.9298 (92.98%)

=====
===== IMPROVED MODEL (With Feature Scaling) =====
=====
Features scaled using StandardScaler

Example of scaling effect on first feature:
Before scaling - Mean: 14.13, Std: 3.52
After scaling  - Mean: -0.00, Std: 1.00

Accuracy with scaling: 0.9649 (96.49%)
```

```
=====
COMPARISON RESULTS
=====
Baseline accuracy (unscaled): 92.98%
Improved accuracy (scaled): 96.49%
Improvement: +3.51% points
Relative improvement: +3.8%
```

```
=====
DETAILED CLASSIFICATION REPORT (With Scaling)
=====
```

	precision	recall	f1-score	support
Malignant	0.97	0.93	0.95	42
Benign	0.96	0.99	0.97	72
accuracy			0.96	114
macro avg	0.97	0.96	0.96	114
weighted avg	0.97	0.96	0.96	114

```
=====
WHY FEATURE SCALING IS CRITICAL FOR KNN
=====
```

K-Nearest Neighbors (KNN) is a distance-based algorithm that calculates the distance between data points to make predictions.

WHY SCALING MATTERS:

1. Features with larger numerical ranges dominate distance calculations
Example:
 - Feature A range: 0-1 (small influence)
 - Feature B range: 0-1000 (huge influence)Without scaling, Feature B would completely dominate the distance!
2. Equal Contribution:
 - Scaling ensures all features contribute equally to the distance calculation
 - No single feature unfairly influences the results
3. Better Performance:
 - As shown above, scaling improved accuracy from 93.0% to 96.5%
 - This is a 3.8% relative improvement!

StandardScaler transforms features to have:

- Mean = 0
- Standard Deviation = 1

This puts all features on the same scale without changing their distribution shape.